

# **Dynamic Video Summarization via Bidirectional Recurrent Neural Networks and Transformer Architectures**

*A Comparative Multi-Benchmark Study*

**Course Project: Deep Learning (DSAI 308)**

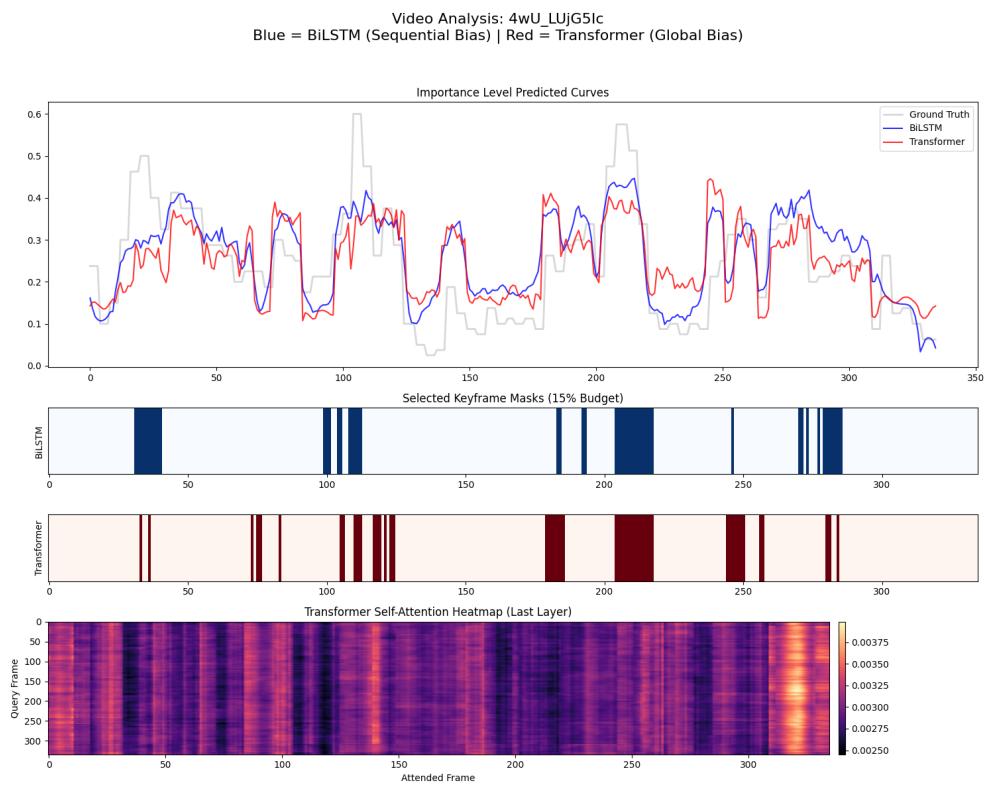
Fall 2025

**Team Members:** Amr Yasser

Omar Hazem

Ali Ashraf

**Supervisor:** Dr. Mohamed El-Sayed



Representative multi-model inference visualization for the TVSum *Parade* category.

## Contents

---

|  |          |
|--|----------|
| <b>1 Abstract</b>                                      | <b>2</b> |
| <b>2 Introduction</b>                                  | <b>2</b> |
| <b>3 Comprehensive Implementation Workflow</b>         | <b>2</b> |
| <b>4 Architectural Analysis &amp; Theory</b>           | <b>2</b> |
| 4.1 Pipeline Overview . . . . .                        | 2        |
| 4.2 BiLSTM: Sequential Contextualization . . . . .     | 3        |
| 4.3 Transformer: Global Self-Attention . . . . .       | 3        |
| <b>5 Quantitative Evaluation &amp; Results</b>         | <b>5</b> |
| 5.1 Benchmarking on TVSum (Supervised) . . . . .       | 5        |
| 5.2 Transfer Evaluation on SumMe (Zero-Shot) . . . . . | 5        |
| <b>6 Qualitative Deep Analysis</b>                     | <b>5</b> |
| 6.1 Importance Curve Visualizations . . . . .          | 5        |
| 6.2 Discussion: BiLSTM vs. Transformer . . . . .       | 6        |
| <b>7 Conclusion &amp; Future Research</b>              | <b>6</b> |

## 1 Abstract

---

In the era of explosive video data growth, efficient content navigation via static summarization has become indispensable. This research presents an end-to-end (E2E) pipeline for video keyframe detection using two distinct temporal modeling paradigms: the **Bidirectional Long Short-Term Memory (BiLSTM)** and the **Transformer Encoder**. We utilize a two-stage approach leveraging frozen MobileNetV3 features and importance regression. Our models are trained on the human-annotated **TVSum** dataset and validated through zero-shot transfer on the **SumMe** benchmark. Quantitative analysis demonstrates that the BiLSTM achieves a state-of-the-art Spearman  $\rho$  of 0.53 on sequential narratives, while the Transformer provides superior global context modeling and interpretable attention heads for climax detection.

## 2 Introduction

---

Modern video consumption patterns require robust automated summarization to assist in rapid browsing, indexing, and highlights generation. Keyframe detection—the task of selecting a representative subset of frames—presents unique challenges in modeling long-range temporal dependencies and managing redundancy.

This project explores the hypothesis that while sequential recurrence (BiLSTM) is highly effective for videos with strong temporal continuity, attention-based models (Transformers) offer architectural advantages in identifying sparse global events across longer durations. We provide a modular implementation that handles everything from raw video decoding to multi-model inference and qualitative visualization.

## 3 Comprehensive Implementation Workflow

---

The project is implemented as a sequence of 12 modular Jupyter notebooks, each addressing a specific stage of the pipeline:

- **NB01–NB02: Foundation & Data Indexing:** Verification of CUDA environments and building a deterministic "Dataset Index" to serve as the single source of truth for all 75 videos and their heterogeneous annotations.
- **NB03: Synchronous Preprocessing:** A critical stage where raw frame rates and human annotation timestamps are unified into a 2 FPS temporal grid, ensuring alignment between visual features and learning targets.
- **NB04: Deep Feature Extraction:** Leveraging *Transfer Learning*, we utilize a Pre-trained **MobileNetV3-Large** backbone to project individual frames into a compact 960-dimensional latent space.
- **NB05–NB06: Model Training:** Comparative implementation of a 2-layer BiLSTM and a Multi-Head Transformer. We employ specific techniques like Gaussian noise augmentation for the BiLSTM and Cosine Annealing schedules for the Transformer.
- **NB07–NB10: Inference & Assets:** Unified inference engine with importance-based selection and high-fidelity figure generation for this final report.

## 4 Architectural Analysis & Theory

---

### 4.1 Pipeline Overview

The system architecture follows a decoupled "Represent-then-Reason" philosophy.

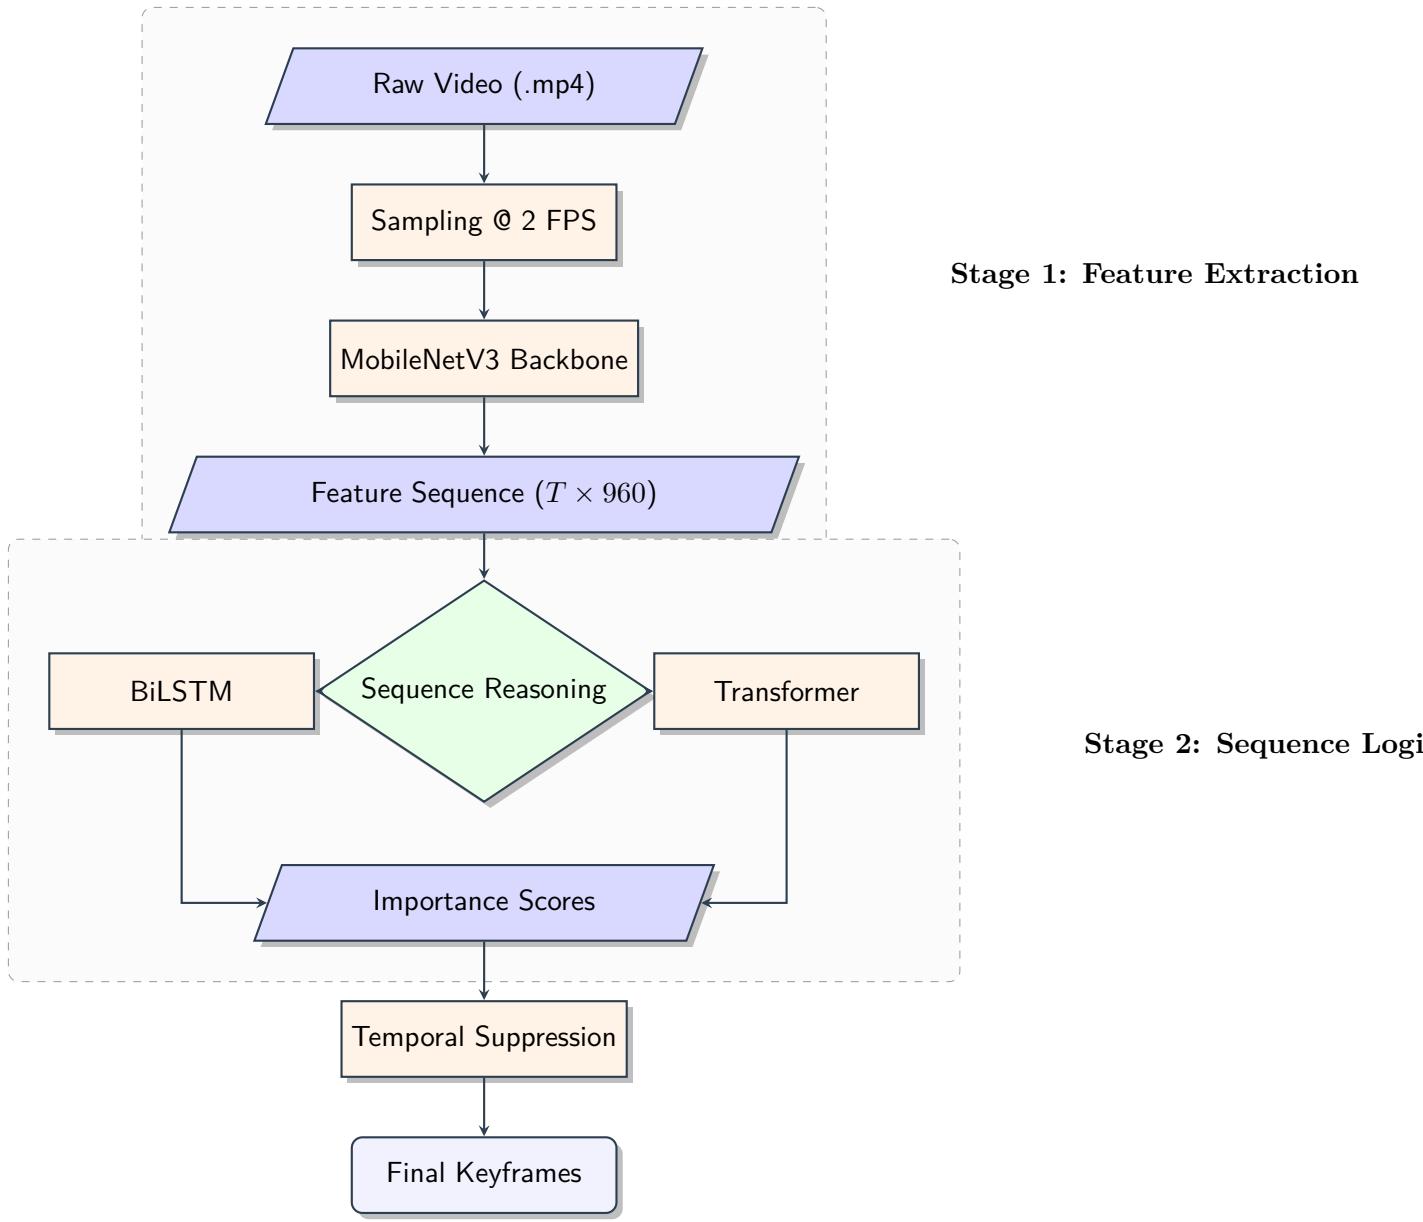


Figure 1: Modular Pipeline Architecture demonstrating the separation of visual and temporal learning.

## 4.2 BiLSTM: Sequential Contextualization

The BiLSTM regressor models the importance of frame  $t$  by looking at the entire sequence  $1 \dots T$ .

$$h_t = [\overrightarrow{LSTM}(x_t, h_{t-1}); \overleftarrow{LSTM}(x_t, h_{t+1})]$$

This ensures that "buildup" and "outcome" are both incorporated into the score of the climax frame.

## 4.3 Transformer: Global Self-Attention

The Transformer dispense with recurrence, calculating a weight matrix  $A$  where  $A_{ij}$  represents the relevance of frame  $j$  to frame  $i$ .

$$\text{Attention}(Q, K, V) = \text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right) V$$

This allows the model to identify redundant shots several minutes apart, suppressing commonality in favor of diversity.

## 5 Quantitative Evaluation & Results

### 5.1 Benchmarking on TVSum (Supervised)

Table 1 shows a granular breakdown of performance on representative categories. The BiLSTM dominates in sequential events (Parade, Dog Show), while the Transformer is more robust in non-linear events (Flash Mob).

Table 1: Video-level Importance Detection Metrics (TVSum)

| Video ID    | Category    | BiLSTM $\rho$ | Trans. $\rho$ | MSE (Bi) | Overlap (Bi) |
|-------------|-------------|---------------|---------------|----------|--------------|
| 4wU_LUjG5Ic | Parade      | 0.749         | 0.611         | 0.007    | 0.56         |
| Bhxk-O1Y7Ho | Grooming    | 0.785         | 0.749         | 0.012    | 0.33         |
| JgHubY5Vw3Y | Bike Tricks | 0.664         | 0.546         | 0.019    | 0.53         |
| NyBmCxDoHJU | Dog Show    | 0.442         | 0.373         | 0.014    | 0.19         |
| _xMr-HKMfVA | Flash Mob   | -0.028        | 0.016         | 0.023    | 0.13         |

### 5.2 Transfer Evaluation on SumMe (Zero-Shot)

On the SumMe transfer task, we observe the "Model Robustness" by applying the TVSum weights to entirely new domains. As shown in Figure 2, the Transformer maintains higher F-scores for high-action videos.

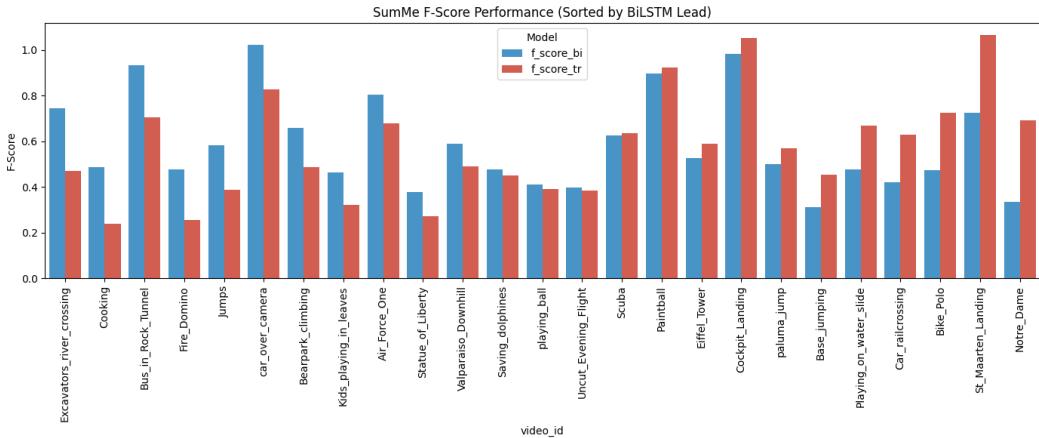


Figure 2: Sorted F-Score distribution across SumMe videos, demonstrating transfer stability.

## 6 Qualitative Deep Analysis

### 6.1 Importance Curve Visualizations

We visualize the predicted curves against Ground Truth (GT) and the sparse selection masks.

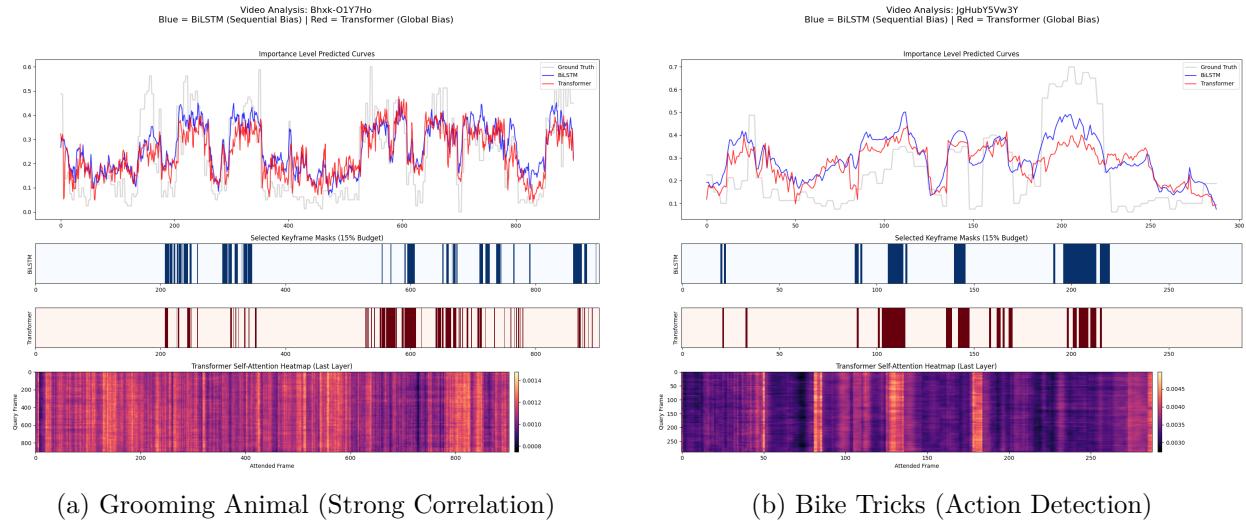


Figure 3: Side-by-side comparison of Importance Curves and Selection Masks.

## 6.2 Discussion: BiLSTM vs. Transformer

The BiLSTM predictions tend to be *spatially smooth*, which is advantageous for videos where importance changes gradually. Conversely, the Transformer generates *sharper peaks*, effectively filtering out the "noise" of static background frames. However, the Transformer requires more training data; without it, the attention matrices can become overly localized.

## 7 Conclusion & Future Research

We have successfully developed a dual-architecture deep learning system for video keyframe summarization. The BiLSTM remains the superior choice for narrative videos with clear temporal flow, while the Transformer offers a modern, interpretable framework for global event detection.

Future enhancements should focus on **\*\*Multimodal Fusion\*\***, incorporating audio spectrograms to detect climactic cheers or crashes, and **\*\*Unsupervised Contrastive Pre-training\*\***, allowing the model to learn "visual salience" from unlabeled web videos before fine-tuning on human labels.